

STIR Overview for developers

Kris Thielemans

version 6.4

SPDX-License-Identifier: Apache-2.0 AND License-ref-PARAPET-license

See STIR/LICENSE.txt for details.

Contents

1	Introduction	2
1.1	Language support	2
2	General conventions	3
2.1	Units	3
2.2	Coordinate system etc	3
2.2.1	Image data	3
2.2.2	Info on projection data	3
3	Code conventions	6
3.1	Configuration file	6
3.2	Namespace	6
3.3	Naming conventions	6
3.4	File conventions	7
3.5	Documentation	7
3.6	Class definitions	8
3.7	Argument order conventions	8
3.8	Error handling	9
3.9	Diagnostics	9
3.10	C++ conventions	9
3.11	Advanced (?) C++ features used	9
3.11.1	Templates	9
3.11.2	Run Time Type Information	9
3.11.3	Iterators	10
3.11.4	Shared pointers and other smart pointers	10
3.12	Generic functionality of STIR classes	11
3.12.1	Copying objects	11
3.12.2	Comparing objects	11
3.12.3	Parsing from text files	12
3.12.4	typedefs	12
4	CUDA support	12
5	Overview of classes	12
5.1	Images	12
5.2	Projection data classes	13
5.3	List-mode data	14
5.4	Data (or image) processor hierarchy	14
5.5	Projector classes	14
5.5.1	Symmetries	15
5.5.2	ForwardProjectorByBin hierarchy	15
5.5.3	BackProjectorByBin hierarchy	15
5.5.4	ProjMatrixByBin hierarchy	16

5.5.5	ProjMatrixElementsForOneBin	17
5.6	Objective functions	17
5.7	Reconstruction classes	17
6	Parsing from text files (or strings)	18
7	IO	19
7.1	Images	19
7.2	Numerical arrays	19
7.3	Projection data	19
7.4	Dynamic or gated data	19
8	Using class hierarchies	20
9	Registries of classes	21
9.1	How does the STIR implementation of registries work?	21
10	Extending STIR with your own files	21
11	Contributing to STIR	23
11.1	Continuous integration testing	23
12	Conclusion	23

1 Introduction

The objective of this document is to describe the STIR common building blocks library. The library has been designed so that it can be used for many different algorithms and scanner geometries, including both cylindrical PET scanners and dual-head rotating coincidence gamma cameras (although the latter needs some work). The library contains classes and functions to run parts of the reconstruction in parallel on distributed memory architectures, but the parallel features are not discussed here.

The building block classes discussed in this document are:

- Classes for images (2D and 3D)
- Classes for projection data (i.e. the measured data)
- Forward projection, back projection and projection matrix classes.
- Classes for iterative reconstruction algorithms
- IO

The documentation for the STIR library is generated automatically from the source using the *doxygen* program (<https://www.doxygen.org/>). This can produce an HTML version but also various other output formats (you can get e.g. RTF or LaTeX output by simply running doxygen on the source files that come in the STIR distribution).

1.1 Language support

STIR is written in C++ and currently requires C++-17, but it is compatible with newer versions of the C++ standard. We will enforce C++-20 from STIR 7.0 (possibly earlier).

Python and MATLAB support is provided via **SWIG**. This means that Python/MATLAB interfaces follow the C++ classes closely, although some differences are required as these languages do not support templates for instance. Naming currently therefore often includes the type,

e.g. `FloatVoxelsOnCartesianGrid` corresponds to C++ `stir::VoxelsOnCartesianGrid<float>`. Python/MATLAB documentation is auto-generated from the C++ documentation. This is not entirely satisfactory, but the best we can do at present. The SWIG interface files are located in `STIR/src/SWIG/*.i`.

2 General conventions

2.1 Units

Distances are in millimetre.

Angles are in radians.

Relative times are in seconds.

2.2 Coordinate system etc

STIR uses a “gantry-based” coordinate system, which is left-handed.

WARNING: This text is correct for PET data. For SPECT, we currently have problems with the image orientation, caused but the fact that SPECT DICOM raw data follow the patient position (not the gantry). Therefore, LPS info etc is currently **incorrect** for SPECT reconstructed images.

2.2.1 Image data

Although STIR is prepared for general images such as blobs on bcc grids (see the online documentation for class **DiscretisedDensity**), currently only voxels on a Cartesian grid are implemented (class **VoxelsOnCartesianGrid**). The coordinate axes for Cartesian grids are chosen as follows.

x-axis : horizontal axis, pointing right when looking from the bed into the gantry

y-axis : vertical axis, pointing downwards

z-axis : the scanner axis, pointing from the gantry towards the bed

The origin of the X and Y axes are located on the central axis of the PET scanner and the Z origin ($z=0$) is located in the middle of the first ring (i.e at the opposite side of the bed). Note that for images with an even size in x and y, the axis of the scanner *does* coincide with the centre of a pixel. In particular, for range of $2n$, the (internal) image coordinates would run from $-n$ to $(n-1)$.

2.2.2 Info on projection data

Conventional axes and units used for projection data are shown in Figures 2 and 1..

See the STIR glossary for some info on naming conventions for projection data.

In 3D PET, two data storage modes are generally used for a **segment**¹

- where the 3D data is ordered by **sinogram** (i.e **axial position**, **view angle**, **tangential position**) (see Figure 3)
- where the 3D dataset is ordered by **view** (i.e **view angle**, **axial position**, **tangential position**) (see Figure 4)

This notation means that for a sinogram, the tangential position runs fastest.

In both modes, the 3D dataset has been stored in several segments where the number of axial positions in each **segment** depends on the **axial compression (span)**.

From both modes, 2 different types of 2D datasets can be obtained :

¹The GE file formats does *not* store the data per segment, but per view. However, once read into STIR, this organisation is no longer available.




Figure 1: Axes and units within one transaxial slice of the target image. The transaxial section of the field of view is shown as a grey circle. The number of measured projection elements along the s -axis is odd, so that $s = 0$ is positioned at the centre of the central projection element. The angles θ and ϕ define the direction of the line-of-response




Figure 2: Sketch of main axes, units and angles used in the cylindrical scanner geometry (axial section, not to scale)

- Sinogram i.e (View, tangential position) for a fixed axial position.
- Viewgram i.e (axial position, tangential position) for a fixed view.

Note that since STIR version 6.0, Time-of-Flight (TOF) is supported. This introduces another index. However, **Sinogram** and **Viewgram** remain 2D objects, and **Segment*** 3D. This will also be the case once we have layers and energy windows. In STIR 5.2, we have therefore introduced new classes **SinogramIndices**, **ViewgramIndices**² and **SegmentIndices**, containing all “other” indices necessary to get at the corresponding data. This means that the recommended coding style is now

```
ViewgramIndices vg_idx(view_num,segment_num);
auto viewgram = proj_data.get_viewgram(vg_idx);
```

In the next version of STIR, we will introduce extra methods to be able to conveniently loop over all viewgrams etc.




Figure 3: By sinogram storage.

For an N ring scanner, if there is *no* axial compression, there are $N - \text{abs}(\text{segment_num})$ sinograms (or axial positions) in each segment. The first sinogram always corresponds to a coincidence in the first ring of the scanner ($\text{ring_num}=0$), where the other ring would obviously be $\text{ring_num}=\text{abs}(\text{ring_diff})=\text{abs}(\text{segment_num})$.

Well ok, this is really true unless you messed around with the data and removed/added some sinograms.

The situation with axial compression is more complicated. Obviously there is more than 1 ring-difference in a segment now. This makes it kind of hard to count how many sinograms/axial_positions you will have and which ring-pairs sit in which segment. You can find some documentation on how we do this in `ProjDataInfoCylindrical.cxx`. See also the Michelogram documentation in the STIR glossary document.

In any case, currently STIR always assumes for each segment that the middle of the scanner corresponds to the middle of the set of sinograms (taking into account the ring difference).

Without view offset, view 0 currently corresponds to a vertical projection. It also assumes that the first detector is at $x = 0$, $y = -R$ (i.e. on the top of the scanner). Finally, in STIR code, the first ring of the scanner is the ring furthest from the bed. This corresponds to how the data are

²Replacing `ViewSegmentNumbers` in previous versions of STIR.




Figure 4: By viewgram storage.

stored in ECAT6,7 and Interfile files. The “highest” TOF bin corresponds to voxels closer to the “first” detector (as encoded for instance in `stir::DetectionPositionPair`).

The Python example `coordinates_vs_bins.py` could be useful to clarify these conventions.

3 Code conventions

3.1 Configuration file

The file `stir/common.h` contains general configuration info and tries to iron out some incompatibilities for certain compilers. If you include any STIR `.h` file, you are guaranteed to have included `stir/common.h` as well, hence there is usually no need to include it explicitly.

3.2 Namespace

The STIR library has its own C++ namespace: `stir`. All symbols are within this namespace. The effect of this is that conflicts with other symbols are impossible (except when somebody else is using the same namespace). STIR also uses sub-namespaces for certain things. This usage will probably be expanded in the future. In most of the code, we use macros such as `START_NAMESPACE_STIR` (originally used for older compilers).

3.3 Naming conventions

Types start with capitals, every word is capitalised, no underscores, e.g. `DiscretisedDensity`. Variables, methods and members are lower case, underscores between different words, e.g. `voxel.size`. Variables, methods and members indicating

- a variable or member used to indicate the number of things starts with *num*, e.g. `num_gates`.
- the number of an item in a sequence end with *num*, e.g. `gate_num`.³
- a relative time (normally with respect to the scan start) end with *rel_time*, e.g. `tracer_injection_rel_time`.

³We are slowly starting to use `_idx` in a few places as “indices” are a more general concept, and avoid confusion between the pre- and postfix use of `num`.

- Pointers to an object called *something_ptr*. Currently there is no naming distinction between shared pointers etc. This is probably a bad idea. For new code, we recommend to use *something_sptr* for shared pointers, and *something_uptr* for unique_ptrs (see section 3.11.4).

In most cases, access to data of a class is via a *get_something()*, *set_something()* pair.

In general, names are descriptive and hence long. We often take longer to decide about the name than to write the actual code. If you write new code, do the same. You will be grateful when you look back at the code a few months later (and not vilified by your successor).

3.4 File conventions

Most classes have their own *.h*, *.inl*, *.txx* and/or *.cxx* files. File names of such classes are simply *ClassName.h* etc. (preserving capitals). In general, looking at the *.h* file should give enough information on what a class/function does.

The *.inl* files contain inline code for functions. The main purpose for this is to keep the *.h* files as short and clean as possible.

The *.txx* files are similar to the *.inl* files. They contain non-inline definitions of template classes (or functions). They should only be included by the corresponding *.cxx* file where the template is then instantiated for any desired types. This makes it easier for a user to instantiate a template for a new type without modifying the original STIR code.⁴

All STIR include files should be included as *#include "stir/maybe_a_subdir/name.h"*

The *doxygen* comments generally occur in the *.h* file, unless they are very extensive. In this case, the *.h* file contains a brief comment, while the *.inl* or *.cxx* file contains the longer description.

To avoid unnecessary interdependencies between *.h* files (and hence unnecessary rebuilds when modifying one *.h* file), we avoid including *.h* files for 'supporting' classes as much as possible. Instead, we only declare the classes needed, e.g. *class Bin*; instead of *#include "stir/Bin.h"*.

Include files should not contain statements like *using namespace std*; or *using std::cos*; (exceptions for the latter are by intention in *textttshared_ptr.h* and *textttformat.h*).

3.5 Documentation

Most documentation is generated via in-code comments using *doxygen*. There are some overview *LaTeX* files in the *documentation* folder, such as this one.

Documentation of a class and its methods generally works as follows

```

/#!/
  \brief 1-2 line description
  \ingroup somegroup

  Long description
*/
class A
{
public:
  //! brief description
  /#!/
    optional longer description
  */
  void method();
};

```

where both styles of "brief" description are acceptable.

⁴The *.txx* extension as used because the Insight Toolkit (ITK) uses it as well. It stands for template C++ presumably.

Similar styles are used for functions, constants and preprocessor defines. There should also be a header in each file assigning it to a doxygen group.

Doxygen groups have to be defined. This is currently done in `src/include/stir/doxygengroups.h`.

Doxygen comments are also parsed to generate the SWIG documentation. This is currently done via `doxy2swig`, although we hope to transition to SWIG's native doxygen support, once its doxygen parser improves.

3.6 Class definitions

Classes should be defined in their own `.h` file (except in implementation classes, which ideally are in the `detail` namespace). They follow generally the following format (ignoring *doxygen* comments).

```
class A
{
public:
    enumerated types and other typedefs
    static member functions
    data members
    constructors
    destructors
    member functions
protected:
    enumerated types and other typedefs
    static member functions
    data members
    constructors
    destructors
    member functions
private:
    enumerated types and other typedefs
    friends if any
    static member functions
    data members
    constructors
    destructors
    member functions
};
```

Inline member functions are defined in the `.inl` file, and are ordered (where possible) such that if inline member function `a()` calls inline member function `b()`, then `b` should be defined first. (Otherwise very few compilers are able to inline the call to `b` from `a`.)

3.7 Argument order conventions

When passing output arguments (by reference or by pointer), those arguments occur FIRST in the list. For instance

```
void some_function(int& output, const SomeType& some_input_argument, ...);
```

This order of arguments allows the use of default arguments. Note that the standard C++ library does not use this convention.

3.8 Error handling

Problem reporting is via two functions *error()* and *warning()*. Currently, *error()* throws an exception. If you do not “catch” that exception, the program will be simply aborted (after writing a diagnostic message).

In many places, validity of input arguments or of the state of an object is checked by *assert* macros. This code is only compiled when the *NDEBUG* preprocessor macro is not defined (for instance, when setting *CMAKE_BUILD_TYPE* to *Debug*), such that a production version of the programs is not slowed down. If an assertion is false, the program aborts with info on the file and line number where the assertion failed.

3.9 Diagnostics

Many function use the *info()* function, which will write to *stdout*. It has a verbosity argument, such that more frequent diagnostic messages can be disabled, see *Verbosity::set_verbosity*. Setting the level to 0 suppresses all diagnostics output. This is a global setting.

3.10 C++ conventions

STIR uses C++-17 and is compatible with C++-20 to the best of our knowledge. For legacy reasons, we have some preprocessor macros (see *stir/common.h*) to define the stir namespace (see below). These should disappear at a later development stage of the library.

We use *override* for virtual functions in derived classes. This is clearer, less error-prone and avoids compiler warnings.

STIR has a very long history. Lots of code was written pre-C++-11. We gradually simplify code when editing. This is definitely visible for code using iterators, which could be simplified by using *auto* as in this guide. It is a long term project to remove all compiler warnings from the code. Please help.

3.11 Advanced (?) C++ features used

This section describes some C++ features used in STIR that might not be so familiar to non-C++ programmers.

3.11.1 Templates

The library uses templates very often. This allows us to write ‘generic’ code, independent of specific types. For instance, multi-dimensional arrays correspond to

```
template <int num_dimensions, typename elemT>
class Array;
```

To avoid linking problems, the templates which we use are explicitly instantiated in the relevant *.cxx* files. If a user needs other types, (s)he will have to add the instantiations. In some cases, this is not necessary as all methods of a template class are inlined, or the actual definition is in a *.txx* file that you can include.

We do use partial class template specialisation in some places. However, (very ugly) work-arounds are provided for compilers that do not support this feature (although these are being gradually removed).

Very occasionally we use member templates.

3.11.2 Run Time Type Information

Another C++ feature that we use is Run Time Type Information (RTTI). We almost exclusively use this to check validity of pointer (or reference) casts down a hierarchy (‘down-casting’). See section 8 for an example.

3.11.3 Iterators

An important C++ concept, used fairly often in STIR, is an *iterator*. The easiest way to think about iterators is as a sort of generalised pointer, used to iterate through a collection of objects of the same type. For instance, the following code would add 2 to all the elements of a vector:

```
void f(vector<int>& v)
{
    for (auto iter = v.begin(); iter != v.end(); ++iter)
        *iter += 2;
}
```

So, just as pointers, you increment an iterator to go to the next element of the vector, and you use **iter* to access the object that the iterator refers to. However, the advantage is that the above code would just as well work for any other type of collection, as long as it provides an iterator interface. So, typical code in C++ would look as follows:

```
template <class Container>
void f(Container& v)
{
    for (auto iter = v.begin(); iter != v.end(); ++iter)
        *iter += 2;
}
```

Iterators are used with great success in the C++ Standard Template Library (STL), and STIR provides iterators for its container classes. In particular, for multi-dimensional arrays, it is possible to iterate through the array in 2 ways: using *Array<n,T>::iterator* whose iterators point to arrays of dimensions *n-1*, or using *Array<n,T>::full_iterator* which essentially provides a one-dimensional look at the whole array. So, adding 2 to all the elements of a multi-dimensional array can be done as follows

```
template <int n, class elemT>
void f(Array<n,elemT>& a)
{
    for (auto iter = a.begin_all();
         iter != a.end_all(); ++iter)
        *iter += 2;
}
```

Note that use of the *begin_all()* and *end_all()* members which return *full_iterator* objects. Of course, the above function is just an illustration, as in STIR this can be done by using:

```
Array<n,elemT> a = ...;
a += 2;
```

3.11.4 Shared pointers and other smart pointers

STIR uses the `shared_ptr` class considerably. Shared pointers are a specific type of 'smart pointers'. These have the advantage that they essentially clean up after you. That is, you do not have to call *delete* on them. They are even exception proof (something which you cannot achieve with an ordinary pointer). Generally, the destructor of the smart pointer will make sure that the object it points to is deleted (when appropriate).

`std::unique_ptr` is a standard smart pointer class which is suitable for pointers to objects where there's only one smart pointer for each object. Unfortunately, a lot of STIR was written pre-C++-11 and we do not use this smart pointer class too much, and generally use `shared_ptr` instead. This is somewhat unfortunate as these two smart pointers are generally quite different concepts.

`shared_ptr` is a wrapper around the `std` (or boost⁵) smart pointer class which is suitable when there are (potentially) more than one pointer pointing to the same object. It keeps a reference count such that the object is (only) deleted when the last shared pointer that references it is destructed.

Caveat: if you modify the object of 1 `shared_ptr`, the change obviously applies to all `shared_ptr`s sharing that object.

Caveat: if you assign an ordinary pointer to a `shared_ptr`, you cannot delete the ordinary pointer anymore (it will be done by the `shared_ptr`). As a consequence, you cannot assign an ordinary pointer twice to a `shared_ptr`. It is thus better to go all the way, and not have any ordinary pointers anymore.

Caveat: do not initialise a `shared_ptr` with a pointer that cannot/should not be deleted. For instance, initialising it with the address of a local variable, or even a reference, will cause a delete on an object that is not allocated on the heap, and so probably crash your program.

It is clear that smart pointers are very useful, but also somewhat dangerous. For this reason, most classes do not return `shared_ptr`s to their members (even if they store one). Instead, they return a pointer to a const object. This prevents a user of a class to inadvertently change members of the class-object. This does involve a memory/performance penalty when the user then needs to create a `shared_ptr` himself. For example, you'll see code like this

```
ProjData proj_data(...);
ProjDataInfo const * proj_data_info_ptr =
    proj_data.get_proj_data_info_ptr();
shared_ptr<ProjDataInfo> new_proj_data_info_sptr(proj_data_info_ptr->clone());
// now do something with new_proj_data_info_sptr
```

We provide the convenience function `stir::is_null_ptr()` to test if a pointer is 'null', i.e. does not point to anything. `is_null_ptr` works with ordinary pointers, `shared_ptr`s and `unique_ptr`s. Use it to avoid that your code depends on what type of (smart) pointer you are using. Note however that since C++11, it is more convenient to use the automatic conversion to `bool`.

3.12 Generic functionality of STIR classes

This is a (very incomplete) section describing some functionality that many STIR classes have in common.

3.12.1 Copying objects

When using pointers or references to objects of an abstract type, you can copy them using `clone()`, see section 3.11.4.

When you only want to create an object of the same characteristics, but without copying the data itself (e.g. images of the same size), use `get.empty_copy()`.

3.12.2 Comparing objects

Many STIR classes implement comparison using the usual `==` and `!=` operators. The implementation of these operators is somewhat tricky in class hierarchies. We try to follow the approach described in *Overriding the C++ Operator==, An approach that uses the Template Method design pattern* By Daniel E. Stevenson and Andrew T. Phillips, Dr. Dobb's Journal, June 2003, <http://www.ddj.com/184405409>.

For some classes, you might want to check only if two objects are of the same type. For instance, if images have the same sizes, origin etc, but not necessarily the same voxel values. This can be done using the member function `has_same_characterics`.

⁵Boost smart pointer are probably no longer supported, and will no longer be from STIR 6.1.

3.12.3 Parsing from text files

Many class-hierarchies allow to construct an object by parsing a text file, which uses an Interfile-like syntax. Examples are given in the User's guide. see section 6.

All these classes (and some others) have a member function `parameter_info()` which returns a string with all parameters of the object.

Note that many STIR classes are not completely constructed by parsing. Usually, it is necessary to call a `set_up` function to make the object usable.

3.12.4 typedefs

Many class hierarchies use the same typedefs. For example:

```
class ProjDataInfo
{
protected:
    typedef ProjDataInfo root_type; // root of hierarchy

};
class ProjDataInfoCylindrical: public ProjDataInfo
{
private:
    typedef ProjDataInfo base_type;
    typedef ProjDataInfoCylindrical self_type;
};
```

4 CUDA support

STIR 6.2 introduced first versions of some CUDA code. The include file `cuda_utilities.h` contains some helpers classes to copy arrays to/from the device. This is still work-in-progress as objects need to be copied to be used by the rest of STIR.

5 Overview of classes

This section provides an overview of the main ingredients of the library. Detailed description of these classes is in the documentation. Below we only mention some general features.

5.1 Images

Iterative algorithms generally assume that the activity density can be discretised in some way. That is, the continuous density can be approximated by having a linear combination of some basis-functions. The reconstruction problem will try to estimate the coefficients λ_{ijk} of the discretised density

$$\sum_{ijk} \lambda_{ijk} b_{ijk}(\hat{x})$$

The base class corresponding to this kind of data is *DiscretisedDensity*.

We assume that the set of basisfunctions can be characterised by 3 indices⁶ (*ijk*) such that *i* runs over a range of integers *i*₁..*i*₂, *j* runs over a similar range that can however depend on *i*, and *k* runs over a similar range that can depend on *i* and *j*. This concept of ranges is embodied in the *IndexRange* class. Multi-dimensional arrays which have such ranges are encoded by the *Array* class. This forms the data structure for the set of coefficients of the basisfunctions, hence *DiscretisedDensity* is derived from the *Array*.

⁶Actually most of the image hierarchy is templated in the number of dimensions. This means it can be used for 2D and 3D images, but also for higher dimensional cases (say temporal information).

In most useful cases, the basisfunctions will be translations of a single function $b(x)$ (although scaling etc could occur depending on ijk). This means that the discretisation has a certain grid, corresponding to the centre of the basisfunctions. This structure is the next level in the image hierarchy. Currently we have the class *DiscretisedDensityOnCartesianGrid* to implement the case where the grid is formed by an orthogonal set of vectors. Another case would be e.g. *DiscretisedDensityOnCylindricalGrid* (for a cylindrical coordinate system), but we have not implemented this yet.

The next level in the hierarchy is then finally the specification of the basis functions themselves. We currently have only voxels and pixels, but another useful case would be to use Kaiser-Bessel functions as the translated basisfunction $b(x)$ (this is called *Blobs* in the literature). This leads us to the example image hierarchy in Figure 5.




Figure 5: Hierarchy of classes for images. Only two of the bottom classes are currently implemented.

Although a Cartesian grid seems the logical choice for images, there are two reasons why one would like to use different types of grid. One important reason would be sampling efficiency. As an example, in 3D one needs less sampling points with a BCC grid while still being able to represent the same spatial frequencies. This becomes particularly important when using more complicated basis functions than voxels, like Kaiser-Bessel functions ('blobs'). In this case, calculations to perform the projections are much more expensive, so having less grid points becomes important. A second reason for changing the image grid is to increase the symmetry of the projection matrix. For example, for cylindrical scanners using a cylindrical grid would maximise the amount of repetition in the projection matrix. In particular, this would allow even for large PET scanners pre-calculation of the 'independent' part of the projection matrix, and loading it completely into memory when doing the reconstruction. This will not only speed up the reconstruction, but also enable more accurate models of the acquisition to be used, resulting in potentially better resolution and noise behaviour. We believe that this hierarchy can accommodate all known cases used for reconstructions in 3D PET.

5.2 Projection data classes

Different manufacturers use different sampling of projection space. As an example, the HiDAC uses 'polar' sampling. Similarly, Single Photon Emission Tomography (SPET) or Computed Tomography (CT) data are very similar to PET. By having projection data classes that allow for different geometries and file formats, our library should be useful for these image modalities as well.

The general class to access projection data is called *ProjData*. As in 3D-PET, the projection data are potentially huge, we do not generally store the whole data in memory. Instead, *ProjData* has methods for getting subsets of the data, i.e. *SegmentByView*, *SegmentBySinogram*, *Viewgram*, *Sinogram*, *RelatedViewgrams* (see section 5.5.1). In addition, the *ProjData* class provides access to a *ProjDataInfo* object (see Figure 6). This object completely describes the geometry of the data.

Different file formats (or potentially other types of projection data) are handled by having derived classes of *ProjData*, which provide specific implementations for the data access methods.




Figure 6: Hierarchy of classes for (geometric) information of the projection data.

This hierarchy is at the moment very simple, but can easily be extended to accommodate different file formats.




Figure 7: File formats for projection data

Currently missing is support for fan-beam data.

5.3 List-mode data

This is supported as well for both PET and SPECT. See the *Listmode* base-class and its hierarchy.

5.4 Data (or image) processor hierarchy

We provide a hierarchy for functions that modify an image, *e.g.* by filtering or thresholding it. The base class is *DataProcessor*, which is templated in **DataT**, to allow using other types, not only images. Aside from the virtual functions that specify the interface to an image processor, it also provides the basics for a registry of all data processors. See section 9 for some details. Note that the registry is specific to every **DataT**.

5.5 Projector classes

The next important type of ingredient for an iterative algorithm are the projection operations. The forward projection models the measurement. Generally, reconstruction algorithms are some kind of inversion procedure for the following problem

$$y_b = \sum_v P_{bv} \lambda_v$$

where y_b are the measured data. The exact interpretation of the projection matrix depends on the algorithm (it is most of the time probabilistic, in the sense that the above equation would only hold 'on average'). However, this need not concern us here.

Many algorithms (in particular EM-type algorithms) can be written solely in terms of multiplication with the projection matrix (*forward projection*) or with its transpose (*back projection*). This is why we have a *ForwardProjectorByBin* and *BackProjectorByBin* hierarchy. The basic objects handled by these projector classes are *RelatedViewgrams* and (a derived class from) *DiscretisedDensity*. Other algorithms (for instance ART or listmode algorithms) need more detailed access to the elements of the projection matrix, and hence will work with a *ProjMatrixByBin* object⁷.

5.5.1 Symmetries

An important feature of geometric projectors is that several parts of the projection matrix are the same. This is because different (generalised) voxels and bins are related by symmetry. It is important to make use of these symmetries for two reasons. It allows storing only a part of the projection matrix (useful for disk storage and caching), and in on-the-fly projection operations it can be used to speed up the computation, as there is no need to recompute the related elements. As the number of projection data elements related by symmetry depends on the specific geometries (e.g. which class derived from *DiscretisedDensity* is used), we need classes (e.g. *RelatedViewgrams*) that store all related data, classes for describing the symmetries (e.g. *DataSymmetriesForViewSegmentNumbers*), and the necessary operations (*SymmetryOperation*).

At the moment, the only symmetries implemented are specific for Cartesian grids. A discussion of these symmetries is presented in [Par4.1], but also in the online documentation.

5.5.2 ForwardProjectorByBin hierarchy

We provide two derived classes of *ForwardProjectorByBin*:

- *ForwardProjectorByBinUsingRayTracing* computes the P_{bv} elements as the length of intersection of one Line Of Response with the voxel. The actual implementation uses a version of Siddon's algorithm, enhanced to use all possible symmetries. This implementation was discussed in [Par4.1].
- *ForwardProjectorByBinUsingProjMatrix* performs forward projection for any *ProjMatrix*. Indeed, all functionality is already provided by the *ProjMatrix* class. The current class only provides the implementation of the *ForwardProjectorByBin* interface, such that algorithms can use any *ProjMatrix* object available.




Figure 8: Forward projectors

This hierarchy uses the registry mechanism discussed in section 9.

5.5.3 BackProjectorByBin hierarchy

This is very similar to the previous section. We provide two derived classes of *BackProjectorByBin*:

- *BackProjectorByBinUsingInterpolation* computes the P_{bv} elements via interpolation between the projection data for the LOR through the centre of the voxel. We have two different interpolation mechanisms (linear, and piece-wise linear) as discussed in [Par4.1].
- *BackProjectorByBinUsingProjMatrix* performs back projection for any *ProjMatrixByBin*.




Figure 9: Back projectors

This hierarchy uses the registry mechanism discussed in section 9.

5.5.4 ProjMatrixByBin hierarchy

This is a base class for row-wise access to the projection matrix P_{bv} . This class provides 2 essential mechanisms aside from the virtual functions that will be used to get the row of the matrix:

- It can cache the elements of the projection matrix. This caching is obviously useful if you can store the part of the matrix that you need in memory (for instance, a slave might not need the whole matrix). However, even if you cannot store the whole matrix, most algorithms need access to a subset of these elements more than once (for instance, OSEM would need them for forward projecting and for back projecting). Caching can be disabled.
- The application of symmetries is provided at base-class level: the derived classes do not have to bother about this, and can concentrate on computing the 'independent' part of P_{bv} .

At the moment, we have the following derived classes:

- **ProjMatrixByBinUsingRayTracing** uses essentially the same (although more flexible) implementation as the on-the-fly **ForwardProjectorByBinUsingRayTracing**, but returns the row of the projection matrix. However, because symmetries are not handled 'in-line', and elements need to be stored, using a **ProjMatrixByBinUsingRayTracing** object for forward projection is in some cases not as efficient as using **ForwardProjectorByBinUsingRayTracing**.
- **ProjMatrixByBinUsingInterpolation** implements the same as **BackProjectorByBinUsingInterpolation**, but is much slower to initialise.
- **ProjMatrixByBinFromFile** handles the case where the projection matrix is stored on disk [not distributed yet]. Because of the support provided by **ProjMatrixByBin**, only the 'independent' part of the projection matrix needs to be stored. Our current implementation does not yet provide a very compact format for storing the elements (although they are of course stored sparsely).
- **ProjMatrixByBinSPECTUB** is for parallel hole SPECT
- **ProjMatrixByBinPinholeSPECTUB** is for multi-pinhole SPECT




Figure 10: Projection matrix classes

This hierarchy uses the registry mechanism discussed in section 9.

⁷Some algorithms would need column-wise access to the projection matrix (i.e. by voxels). We did not implement any of these algorithms, so we do not have appropriate projection classes for them. Modification of the row-wise classes is straightforward though.

5.5.5 ProjMatrixElementsForOneBin

This class provides sparse storage for a row of the projection matrix. It has methods to access the data using an STL-style iterator (which is essentially a generalised pointer). This means that the actual way to store the data is hidden from the user. In principle, this format could be very compact, or alternatively very efficient. At the moment, we provide a format that is somewhere in between. For instance, image indices are not stored incrementally, as although it would allow very compact storage, it is detrimental for speed.

5.6 Objective functions

Many iterative (reconstruction) algorithms can be formulated in terms of an objective function, which is the algorithms tries to maximise (or minimise). Since STIR 2.0, this is implemented as a class hierarchy based on `GeneralisedObjectiveFunction`, where we use the convention that the objective function is maximised.

Some iterative algorithms use an 'objective function' only in a loose sense. They might for instance allow generalisations which no longer optimise a function. For example in the case of EMLL with non-matching forward and back projectors, the 'gradient' that is computed is generally not the gradient of the log-likelihood that corresponds to the forward projector. However, one hopes that it still points towards the optimum. The corresponding objective function is implemented in the classes

`PoissonLogLikelihoodWithLinearModelForMeanAndProjData` for projection data and `PoissonLogLikelihoodWithLinearModelForListModeData` for list-mode data. There are classes for dynamic and gated data as well.

There can be different objective function that use common operations. For instance, the objective function could implement a least squares criterion, or a Poisson log-likelihood. Both of these have a model for the mean of the measured data (given an image). Generally speaking, STIR implements those common operations using a separate class. As an example, projection operations are implemented via a pointer to a `ProjectorByBinPair`, which in itself is a small class hierarchy using either a `ProjMatrixByBin` object, or a `ForwardProjectorByBin` and `BackProjectorByBin` pair.

Note that classes exist for Poisson log-likelihood for both (histogrammed) projection data and list-mode data.

Often, one includes a penalty (or prior) in the objective function (see doxygen documentation for the class `GeneralisedPrior`). The penalty is expected to be a function that increases with higher penalty, so it will be *subtracted* from the unregularised case.

See the doxygen documentation `GeneralisedObjectiveFunction` for more details.

5.7 Reconstruction classes

The final ingredient for performing reconstructions is the reconstruction algorithm itself. We also have a hierarchy for this, as many iterative algorithms are similar, and differ only in some minor steps. We do not discuss the actual 'leaves' of the Reconstruction tree here.

The tree starts at the abstract `Reconstruction` base class. Classes ending in "*Reconstruction*" are abstract base classes for a family of algorithms. Classes including the name of an algorithm are concrete ones whose principal purpose is to implement the reconstruction method of the algorithm.

Certain algorithms will only work with particular objective functions, while others (such as gradient ascent) might work with arbitrary objective functions.

This arrangement gives the users of the library a flexible range of implementation choices, allowing them to experiment with new algorithms and new variations of old ones with minimal coding effort.

The `Reconstruction` hierarchy is (since STIR 2.0) templated in `TargetT`. This is the type of the output, i.e. normally the image type. Templating this type gives flexibility to have different output-types, such as parametric images or even normalisation factors.

6 Parsing from text files (or strings)

Many STIR classes are based on the `ParsingObject` class, and hence are able to

- ask the parameters to the user interactively
- read the parameters from an Interfile-style file
- return parameter info such that an Interfile-style input file can be created from the current set of parameters.

In the simplest case, the only thing you need to do is to add the variables and keywords to the keymap. All functionality will then follow. See the following example.

```
class A : public ParsingObject
{
    private:
        typedef ParsingObject base_type;

        int number; float a;

        void initialise_keymap() override
        {
            base_type::initialise_keymap();
            this->parser.add_start_key("My Class Parameters");
            this->parser.add_key("number", &number);
            this->parser.add_key("width", &a);
            this->parser.add_stop_key("END My Class Parameters");
        }

        void set_defaults() override
        {
            base_type::set_defaults();
            // specify defaults for the parameters in case they are not set.
            number = 1;
            a = 2.4F;
        }

        bool post_processing() override // will be renamed to post_parsing()
        {
            // do some checks and handle some extra variables
            return false; // everything was ok
        }
};

int main(int argc, char **argv)
{
    A a;
    // parse a file (first argument on the command line)
    a.parse(argv[1]);
    // list them back to cout
    std::cout << a.parameter_info();
    return EXIT_SUCCESS;
}
```

Check the doxygen documentation for `ParsingObject`.

7 IO

7.1 Images

For images, output file formats and (since version 2.0) input file format are represented by objects as well, as instantiations of classes derived from `OutputFileFormat` and `InputFileFormat` respectively. All of the file formats supported by STIR are stored in registries (see section 9).

Currently recommended way to read an image is as follows:

```
typedef DiscretisedDensity<3,float> DataType ;
std::unique_ptr<DataType> density_aptr = read_from_file<DataType>(filename);
```

You will have to include `stir/IO/read_from_file.h` for this to work. The function `read_from_file` will open the file, read an initial block of data, check if that corresponds to the signature of any of the file formats in the default registry, and use the first matching file format to read the data.

For output, you can use

```
write_to_file(filename, density);
```

You will have to include `stir/IO/write_to_file.h` for this to work. This will write using the default output file format and is equivalent to the following:

```
auto output_format_sptr =
    OutputFileFormat<DataType >::default_sptr();
output_format_sptr->write_to_file(filename, density);
```

You will have to include `stir/IO/OutputFileFormat.h` for this to work. Currently, the default output file format in STIR is Interfile (see the User's Guide).

At present, STIR does not provide a way to open an image file for reading and writing. Patient/study information etc is currently ignored. More dangerously, image orientation as well.

It is relatively straightforward to add a new file format by deriving a new class from `InputFileFormat` (for input), and adding that to the default registry (see section 9 for more details).

7.2 Numerical arrays

There are some functions to read/write numerical vectors/arrays from file. Check `stir/IO/read_data.h` (and `write_data.h`) for binary IO, and `stir/stream.h` for text-based IO.

7.3 Projection data

At present, there is no registry yet for file formats of projection data. For reading, use

```
shared_ptr<ProjData> proj_data_sptr = ProjData::read_from_file(filename);
```

This has an optional argument for read/write access to the file.

For writing, you could use

```
ProjDataInterfile proj_data(proj_data_info_sptr, filename);
```

We currently have no easy way to copy all data across from one `ProjData` object to another. You will have to write an explicit loop over segments.

7.4 Dynamic or gated data

Data should have an `ExamInfo` with a `TimeFrameDefinitions` object. You can also use the `get_time.frame.info` utility. There is no equivalent yet for gating fractions.

8 Using class hierarchies

The main strength of object-oriented programming is that you can write code that is as generic as possible⁸. In STIR, we have tried to use this feature as much as possible (although there are still remaining generalisations to be implemented). This works by having a base class and a bunch of derived classes. Examples of class hierarchies are given above. The question is however how you use these hierarchies. Some attempt is given here to answer this, but of course, for serious work, you really should read a good C++ book.

We will use the `DiscretisedDensity` hierarchy as an example, but the same holds for nearly all other hierarchies in STIR (or indeed C++). In addition, you will see some templates being used.

Rule 1: You should try to write your code using the type of the base-class. If this is not possible, use a class as close to the base-class as you can.

Example: If you have a function that manipulates images, write it in terms of `DiscretisedDensity`, not of `VoxelsOnCartesianGrid`.

Rule 2: Pass references or smart pointers as arguments of functions. Use references if you can (such that your code would work for any type of pointer). Never (or almost never) return references, but return smart pointers (or a const pointer to an object that is guaranteed to stay alive long enough).

Example:

```
template <int num_dimensions, typename elemT>
shared_ptr< DiscretisedDensity<num_dimensions, elemT> >
do_something(const DiscretisedDensity<num_dimensions, elemT>& density1,
             const DiscretisedDensity<num_dimensions, elemT>& density2);
```

This is necessary, as you generally cannot construct an object of the type of the base-class. Similarly, it is impossible to copy an object of a base-class. This is solved by using the *virtual copy constructor* idiom:

```
shared_ptr< DiscretisedDensity<3,float> > new_density_sptr(density.clone());
```

This constructs an object of the same type as the given density.

Of course, many functions do need to know the actual type of the parameters, either because they can work with only one type (for instance, `ForwardProjectorByBinUsingRayTracing` handles only `VoxelsOnCartesianGrid`), or because the processing will be different for different types. This can be solved by using C++ Run Time Type Information (RTTI), preferably in the form of `dynamic_cast`.

```
template <typename elemT>
void f(DiscretisedDensity<3, elemT>& density)
{
    auto voxels_ptr =
        dynamic_cast< VoxelsOnCartesianGrid<elemT> * > (&density);
    if (!voxels_ptr)
        error("f: can only handle images of type VoxelsOnCartesianGrid");
    ...
}
```

Note that `dynamic_cast` of pointers allows easy checking if the types matched. When using references, a non-matching type will cause an exception to be thrown. There currently is no code in STIR using/catching exceptions, so if you do not catch the exception yourself, your program will abort mysteriously.

⁸In C++, another very powerful way to write generic code are templates, see section 3.11.

9 Registries of classes

Several class hierarchies in STIR keep a registry of all *leaf* derived classes (i.e. classes at the end of the hierarchy). This registry allows the user to select *at run-time* which image processor, projector, matrix, output file format etc. (s)he wants to use. It is fairly easy to add your own leaf class to the hierarchy such that it will be added to the registry. How to do this is briefly described here, but see the doxygen documentation of classes `RegisteredObject` and `RegisteredParsingObject` and the example below.

The recommended way to add your own class (as always) is to take a look at an existing leaf in the hierarchy, copy the files and change what you need. Also make sure your new class gets linked in by adding it to one of existing `*registries.cxx` files, although it is probably better to put your own additions in separate files such that you can easily import a new version of STIR. See section 10 for information on how to do that.

9.1 How does the STIR implementation of registries work?

The `RegisteredObject` class takes care of finding your class etc, but it will only do that if you have registered your own class. In an attempt to avoid having to do this explicitly, we currently use the following scheme:

- creating a variable of type `yourClass::RegisterIt` will add *yourClass* to the registry as long as this variable persists.
- we put that variable in a separate source file and declare it static. This means it will be “initialised” (*i.e.* its constructor will be called, which actually does the registration) before the main program starts.
- the special source file needs to be added explicitly to the list of files to link with⁹. This is currently done by adding the file to the `STIR_REGISTRIES` CMake variable.

10 Extending STIR with your own files

(See the section on using STIR in an external project in the User’s Guide if you just want to use STIR functionality in your own work.)

If you want to let STIR use your own extra classes, *e.g.* if you added your own filter or prior and want to use that during reconstruction, you can do this while keeping your own work independent of the STIR sources.

- Put all your files in a separate folder (which can be outside the STIR tree) or in `STIR/src/local` (the distribution will never contain any files in that directory). Let us call your directory **mySTIRext** as an example.
- Create a file `mySTIRext/extra_stir_dirs.cmake` with CMake statements, in particular a list of your subdirectories. For instance,

```
# add my include directory to compiler switches (or better: use target_include_directories)
include_directories(include)

# check CMakeLists in next directory
# note: 2nd argument needs to have different name than any of the existing STIR libraries
add_subdirectory(${STIR_LOCAL}/buildblock my_buildblock)

# add registries (if you have any)
```

⁹If the file is added to a library, many linkers will think that it only contains variables which are not referenced by anything else, so will discard it from the linking, and hence the relevant constructor will never be called.

```
list(APPEND STIR_REGISTRIES
    ${STIR_LOCAL}/buildblock/my_registries.cxx)

# add to list of libraries to include in linking
list(APPEND STIR_LIBRARIES my_buildblock)

# copy variables to PARENT_SCOPE
SET( STIR_REGISTRIES ${STIR_REGISTRIES} PARENT_SCOPE)
SET( STIR_LIBRARIES ${STIR_LIBRARIES} PARENT_SCOPE)
```

where **buildblock** is the name of your subdirectory and **my_buildblock** is a name for the sub-directory where CMake will build the files (make sure the latter is different from any sub-directories in **STIR/src**).

- Put in each of your subdirectories a file called **CMakeLists.txt** which list all files you want to compile. This file have to be in a specific (but simple) format. The easiest way to do this is to copy for instance **buildblock/CMakeLists.txt**, **utilities/CMakeLists.txt** or **test/CMakeLists.txt** and change the list of filenames. For instance, an extra library would be specified like this:

```
set(dir my_buildblock)
set (dir_LIB_SOURCES ${dir}_LIB_SOURCES)

set(${dir_LIB_SOURCES}
    my_source1.cxx
    my_source2.cxx
)

include(stir_lib_target)

# Declare dependencies on other STIR libraries
target_link_libraries(${dir} PUBLIC ${STIR_LIBRARIES})
# Alternatively, only depend on a specific STIR libraries (NOT recommended as their names will
# For instance to depend only on the "normal" STIR buildblock
target_link_libraries(${dir} PUBLIC ${STIR_BUILDBLOCK_LIB})
```

- Re-run **CMake** and set the CMake variable **STIR_LOCAL** to the directory where your **extra_stir_dirs.cmake** file (*i.e.* **mySTIRext** in this example), the result will be that, when you (re-)build STIR, your files will be included in the normal build/test process.

As discussed in section 9, if you implement a projector, filter or similar, you will need to add it to its relevant registry such that normal STIR utilities can use it at run-time. As an example, let us say you have written a class that implements a different output fileformat. You could create a new file **my_registry.cxx** with content like this

```
#include "myOutputFileFormat.h"
static myOutputFileFormat::RegisterIt dummy1;
```

If you add that to the **STIR_REGISTRIES** variable in your **extra_stir_dirs.cmake** as above, everything should work as expected.

WARNING:

Until we switch to C++-20, we use the **fmt** library to provide functionality similar to **std::format**. This is done using a git submodule and a header-only dependency. Please do **not** use **stir/format.h** in your header files (.h or .inl), as this will create a dependency on **fmt** include files for external programs, more than likely leading to compilation problems.

11 Contributing to STIR

See `STIR/CONTRIBUTING.md` and `documentation/devel/README.md` for more information on editor settings etc.

11.1 Continuous integration testing

We use GitHub Actions and Appveyor for testing of every pull-request that was submitted on GitHub. PRs need to pass all checks before they will be accepted of course. To help debugging, we upload `recon_test_pack` log files as artefacts on GitHub if a job fails. You can see these in the Summary of every job that was triggered.

12 Conclusion

In many cases, a reconstruction algorithm does not have to know about the type of the image it gets. In our set-up this means it can/should be implemented in terms of a *DiscretisedDensity* object, and *ForwardProjectorByBin* and *BackProjectorByBin* objects (or *ProjMatrixByBin* object). This means that once for instance EM is implemented, it can be used on voxels, blobs, different grids etc. without having to rewrite the algorithm itself (of course, the projectors do depend on the actual type of image, so you would have to write those).

The STIR library allows maximal code re-use when implementing new functionality. Its structure and its documentation features make it into a unique resource for the PET/SPECT community.